

# Insurance Dataset

## Linear Regression

Problem Statement: To find how charges are varying based on the selected features.

```
In [1]: 1 # importing the necessary libraries
        2 import numpy as np
        3 import pandas as pd
        4 import matplotlib.pyplot as plt
        5 import seaborn as sns
```

In [2]:

```

1 # Data Collection-reading the file
2 df=pd.read_csv(r"C:\Users\yoshitha lakshmi\OneDrive\Desktop\python\insurance.csv")
3 df

```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

In [3]:

```

1 # Data cleaning and Preprocessing
2 df.head()

```

Out[3]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [4]: 1 df.tail()

Out[4]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [5]: 1 df.describe()

Out[5]:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

In [6]: 1 df.shape

Out[6]: (1338, 7)

In [7]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [8]:

```
1 df['sex'].value_counts()
```

```
Out[8]: sex
male      676
female    662
Name: count, dtype: int64
```

```
In [9]: 1 convert={"sex":{"female":1,"male":2}}
        2 df=df.replace(convert)
        3 df
```

Out[9]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.92400
1	18	2	33.770	1	no	southeast	1725.55230
2	28	2	33.000	3	no	southeast	4449.46200
3	33	2	22.705	0	no	northwest	21984.47061
4	32	2	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	2	30.970	3	no	northwest	10600.54830
1334	18	1	31.920	0	no	northeast	2205.98080
1335	18	1	36.850	0	no	southeast	1629.83350
1336	21	1	25.800	0	no	southwest	2007.94500
1337	61	1	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [10]: 1 convert={'smoker':{'yes':1,'no':2}}
          2 df=df.replace(convert)
          3 df
```

Out[10]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.92400
1	18	2	33.770	1	2	southeast	1725.55230
2	28	2	33.000	3	2	southeast	4449.46200
3	33	2	22.705	0	2	northwest	21984.47061
4	32	2	28.880	0	2	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	2	30.970	3	2	northwest	10600.54830
1334	18	1	31.920	0	2	northeast	2205.98080
1335	18	1	36.850	0	2	southeast	1629.83350
1336	21	1	25.800	0	2	southwest	2007.94500
1337	61	1	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

```
In [11]: 1 df=df.drop('region',axis=1)
```

```
In [12]: 1 df=df.drop('children',axis=1)
```

```
In [13]: 1 features=df.columns[0:2]
```

```
In [14]: 1 target=df.columns[-1]
```

```
In [15]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LinearRegression
```

```
In [16]: 1 x=np.array(df[features])
2 y=np.array(df[target])
```

```
In [17]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.7)
2 regr=LinearRegression()
3 regr.fit(x_train,y_train)
4 print(regr.score(x_train,y_train))
```

0.10934585685325382

```
In [18]: 1 print(regr.intercept_)
```

317.7620065058163

```
In [19]: 1 coeff_df=pd.DataFrame(regr.coef_)
2 coeff_df
```

Out[19]:

	0
0	255.553361
1	1178.226540

## Conclusion

The accuracy for this Dataset is very low while using LinearRegression. Accuracy = 0.0835

# Ridge Regression

```
In [20]: 1 from sklearn.linear_model import Ridge,Lasso,RidgeCV
```

```
In [21]: 1 ridgeReg = Ridge(alpha=10)
2 ridgeReg.fit(x_train,y_train)
3 train_score_ridge = ridgeReg.score(x_train,y_train)
4 test_score_ridge = ridgeReg.score(x_test,y_test)
5
6 print('\nRidge Model\n')
7 print('Train score for ridge model is {}'.format(train_score_ridge))
8 print('Test score for ridge model is{}'.format(test_score_ridge))
```

Ridge Model


Train score for ridge model is 0.10932328710163086

Test score for ridge model is0.06876009932822547



In [22]:

```
1 plt.figure(figsize = (10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Ridge;$\alpha=0.7$')
3 plt.plot(features,regr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regr')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```





## Conclusion

For Ridge Regression also the Accuracy value is very low.

Train score for ridge model is 0.08393600452093364

Test score for ridge model is 0.09723547364784435

## Logostic Regression

Problem Statement: To find smokers count based on the features - sex, age.

```
In [23]: 1 import numpy as np
          2 import pandas as pd
          3 import matplotlib.pyplot as plt
          4 import seaborn as sns
          5 from sklearn.model_selection import train_test_split
          6 from sklearn.linear_model import LogisticRegression
```

```
In [24]: 1 df=pd.read_csv(r"C:\Users\yoshitha lakshmi\OneDrive\Desktop\python\insurance.csv")
          2 df
```

Out[24]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [25]: 1 features=df.columns[0:3]
```

```
In [26]: 1 target=df.columns[-3]
```

```
In [27]: 1 from sklearn.datasets import load_digits
          2 digits=load_digits()
```

```
In [28]: 1 print(digits.data.shape)
          2 print(digits.target.shape)
```

```
(1797, 64)
(1797,)
```

```
In [29]: 1 x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test_size=0.7,random_state=2)
          2 lor=LogisticRegression(max_iter=10000)
          3 lor.fit(x_train,y_train)
```

```
Out[29]: LogisticRegression
          LogisticRegression(max_iter=10000)
```

```
In [30]: 1 score=lor.score(x_test,y_test)
          2 print(score)
```

```
0.9523052464228935
```

## Conclusion

Using Logistic Regression score is little high compared to Linear Regression. So, further process continued in Logistic Regression. Accuracy = 0.9523052464228935

## Decision Tree

```
In [31]: 1 import numpy as np
          2 import pandas as pd
          3 import seaborn as sns
          4 import matplotlib.pyplot as plt
          5 from sklearn.model_selection import train_test_split
          6 from sklearn.tree import DecisionTreeClassifier
```

```
In [32]: 1 df=pd.read_csv(r"C:\Users\yoshitha lakshmi\OneDrive\Desktop\python\insurance.csv")
          2 df
```

Out[32]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [33]: 1 df['sex'].value_counts()
```

```
Out[33]: sex
male      676
female    662
Name: count, dtype: int64
```

```
In [34]: 1 df['bmi'].value_counts()
```

```
Out[34]: bmi
32.300    13
28.310     9
30.495     8
30.875     8
31.350     8
..
46.200     1
23.800     1
44.770     1
32.120     1
30.970     1
Name: count, Length: 548, dtype: int64
```

```
In [35]: 1 convert={'sex':{'female':0,'male':1}}
          2 df=df.replace(convert)
          3 df
```

Out[35]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	yes	southwest	16884.92400
1	18	1	33.770	1	no	southeast	1725.55230
2	28	1	33.000	3	no	southeast	4449.46200
3	33	1	22.705	0	no	northwest	21984.47061
4	32	1	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	1	30.970	3	no	northwest	10600.54830
1334	18	0	31.920	0	no	northeast	2205.98080
1335	18	0	36.850	0	no	southeast	1629.83350
1336	21	0	25.800	0	no	southwest	2007.94500
1337	61	0	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [36]: 1 X=['age','sex']
          2 y=['yes','no']
          3 all_inputs=df[X]
          4 all_classes=df['smoker']
```

```
In [37]: 1 X_train,x_test,y_train,y_test=train_test_split(all_inputs,all_classes,test_size=0.7)
```



```
In [38]: 1 clf=DecisionTreeClassifier(random_state=0)
```

```
In [39]: 1 clf.fit(X_train,y_train)
```

```
Out[39]: ▾      DecisionTreeClassifier  
DecisionTreeClassifier(random_state=0)
```

```
In [40]: 1 score=clf.score(X_train,y_train)  
2 print(score)
```

0.7930174563591023

## Random Forest

```
In [41]: 1 import numpy as np  
2 import pandas as pd  
3 import matplotlib.pyplot as plt,seaborn as sns  
4 from sklearn.model_selection import train_test_split
```

```
In [42]: 1 df=pd.read_csv(r"C:\Users\yoshitha lakshmi\OneDrive\Desktop\python\insurance.csv")
          2 df
```

Out[42]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

In [43]:

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [44]:

```
1 x=df.drop('smoker',axis=1)
2 y=df['smoker']
```

```
In [45]: 1 convert={'sex':{'female':0,'male':1}}
          2 df=df.replace(convert)
          3 df
```

Out[45]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	yes	southwest	16884.92400
1	18	1	33.770	1	no	southeast	1725.55230
2	28	1	33.000	3	no	southeast	4449.46200
3	33	1	22.705	0	no	northwest	21984.47061
4	32	1	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	1	30.970	3	no	northwest	10600.54830
1334	18	0	31.920	0	no	northeast	2205.98080
1335	18	0	36.850	0	no	southeast	1629.83350
1336	21	0	25.800	0	no	southwest	2007.94500
1337	61	0	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [46]: 1 from sklearn.ensemble import RandomForestClassifier
          2 rfc=RandomForestClassifier()
          3 rfc.fit(X_train,y_train)
```

Out[46]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [47]: 1 score=rfc.score(x_test,y_test)
          2 print(score)
```

0.7385272145144077

```
In [48]: 1 params={'max_depth':[2,3,5,10,20],
          2         'min_samples_leaf':[5,10,20,50,100,200],
          3         'n_estimators':[10,25,30,50,100,200]}
```

```
In [49]: 1 from sklearn.model_selection import GridSearchCV
          2 grid_search=GridSearchCV(estimator=rfc,param_grid=params,cv=2,scoring="accuracy")
          3 grid_search.fit(X_train,y_train)
```

```
Out[49]:  ▶ GridSearchCV
          ▶ estimator: RandomForestClassifier
            ▶ RandomForestClassifier
```

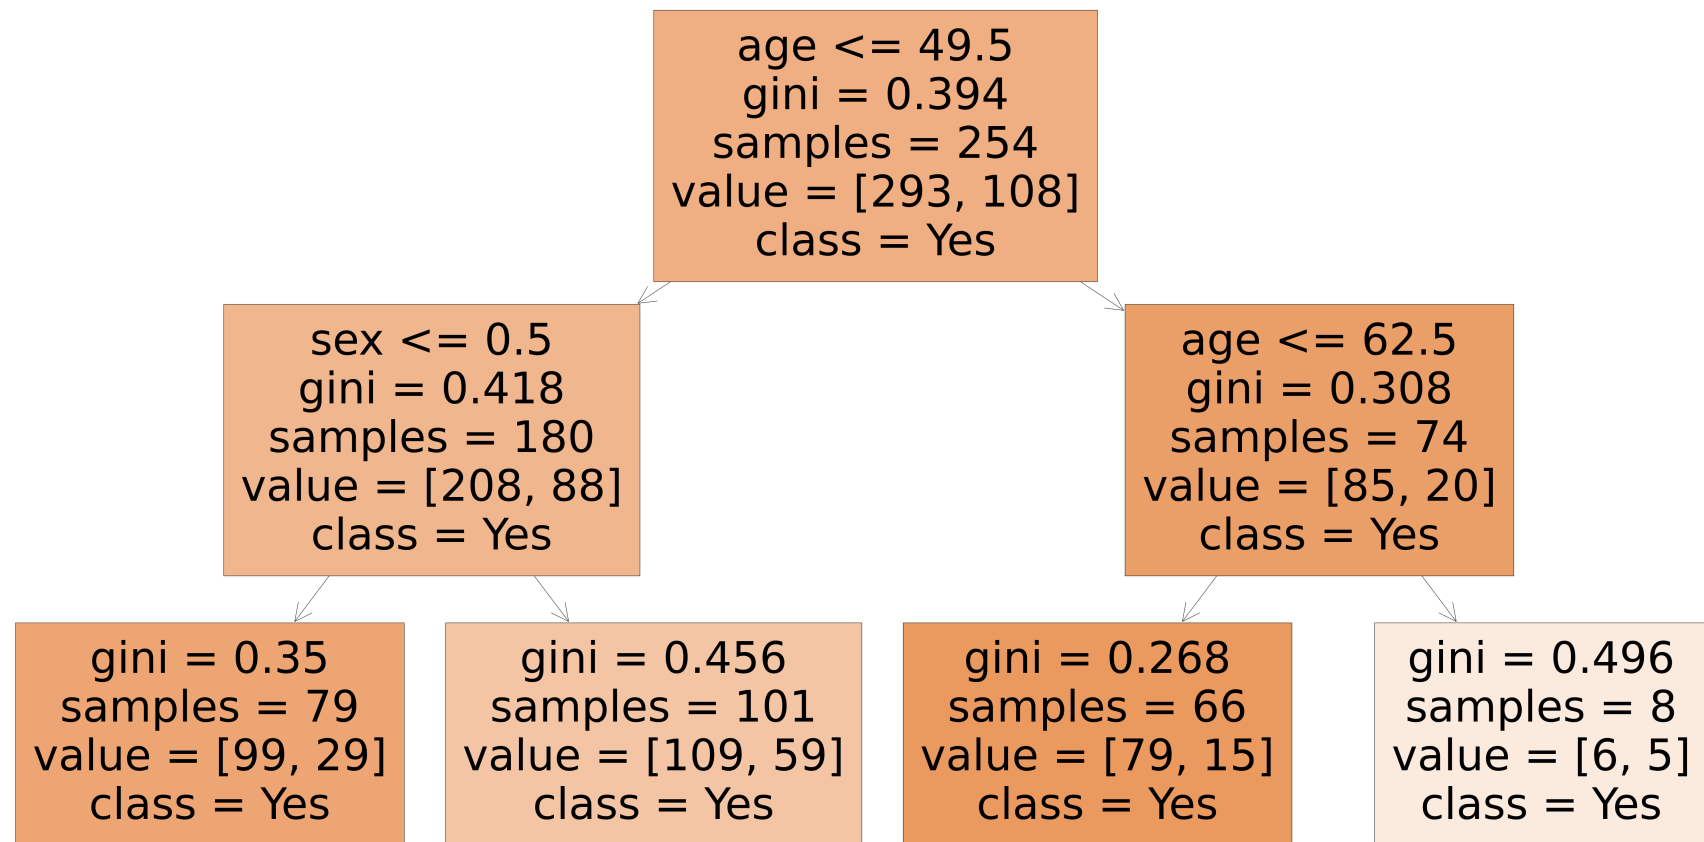
```
In [50]: 1 grid_search.best_score_
```

Out[50]: 0.7780597014925373

```
In [51]: 1 rf_best=grid_search.best_estimator_
```

```
In [52]: 1 from sklearn.tree import plot_tree
2 from sklearn.tree import DecisionTreeClassifier
3 plt.figure(figsize=(80,40))
4 plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[52]: [Text(0.5, 0.8333333333333334, 'age <= 49.5\nngini = 0.394\nsamples = 254\nvalue = [293, 108]\nnclass = Yes'),
Text(0.25, 0.5, 'sex <= 0.5\nngini = 0.418\nsamples = 180\nvalue = [208, 88]\nnclass = Yes'),
Text(0.125, 0.16666666666666666, 'gini = 0.35\nsamples = 79\nvalue = [99, 29]\nnclass = Yes'),
Text(0.375, 0.16666666666666666, 'gini = 0.456\nsamples = 101\nvalue = [109, 59]\nnclass = Yes'),
Text(0.75, 0.5, 'age <= 62.5\nngini = 0.308\nsamples = 74\nvalue = [85, 20]\nnclass = Yes'),
Text(0.625, 0.16666666666666666, 'gini = 0.268\nsamples = 66\nvalue = [79, 15]\nnclass = Yes'),
Text(0.875, 0.16666666666666666, 'gini = 0.496\nsamples = 8\nvalue = [6, 5]\nnclass = Yes')]
```



```
In [53]: 1 imp_df=pd.DataFrame({"varname":X_train.columns,"Imp":rf_best.feature_importances_})
          2 imp_df.sort_values(by="Imp",ascending=False)
```

Out[53]:

	varname	Imp
0	age	0.798067
1	sex	0.201933

## Coclusion

For both Decision Tree and Random Forest

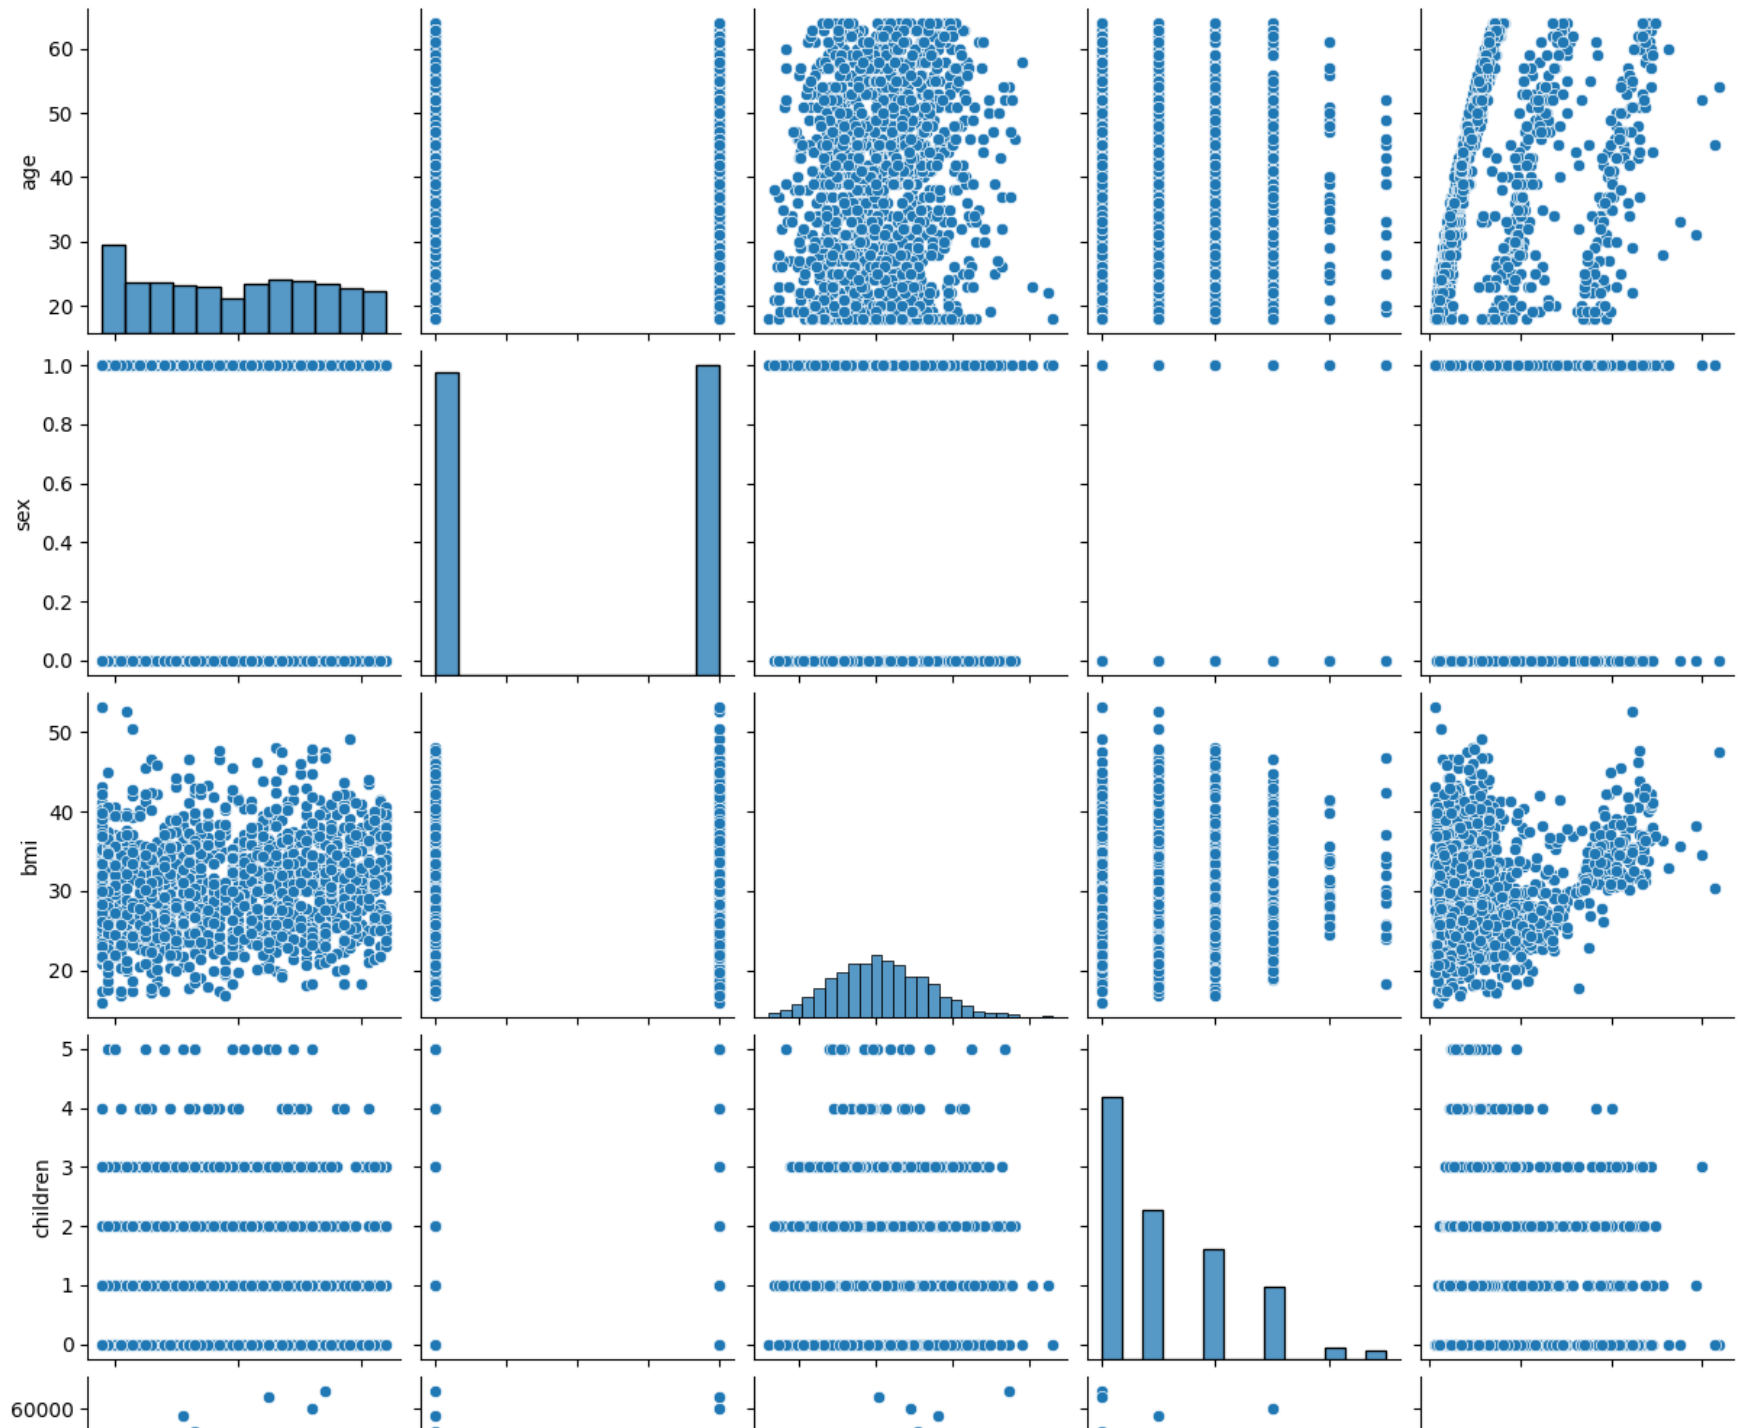
Based on our Problem Statement we classified data and build using Random Forest.


## Exploratory Data Analysis

```
In [54]: 1 sns.pairplot(df)
```

```
Out[54]: <seaborn.axisgrid.PairGrid at 0x1caa19f5420>
```









```
In [55]: 1 sns.displot(df['smoker'])
```

```
Out[55]: <seaborn.axisgrid.FacetGrid at 0x1cabaa07c40>
```



```
In [56]: 1 sns.displot(df['charges'])
```

```
Out[56]: <seaborn.axisgrid.FacetGrid at 0x1cabaa953c0>
```



## Conclusion

Using Exploratory Analysis, the relation between features has discovered.

In [57]: 1 **import** pickle

In [58]: 1 df="prediction"  
2 pickle.dump(rfc,open(df,'wb'))

In [ ]: 1